


Methods (Qwen3.5-2B, Cochrane rephrase-only test split, N=667)

Method	Description
Identity	Reference baseline that copies the input sentence unchanged (no model). Bounds SARI from above and exposes its bias toward light edits.
Baseline	Raw Qwen3.5-2B, zero-shot <code>default_zero_shot</code> prompt, greedy decoding. No retrieval, no fine-tuning.
RAG	Raw Qwen3.5-2B with the <code>definition_augmented</code> prompt: an exact-match biomedical glossary (MedSimplify) injects term definitions into the prompt (84.3% of sentences matched, avg. 2.51 terms).
QLoRA	Qwen3.5-2B fine-tuned with QLoRA on the Cochrane rephrase-only train split (<code>default_zero_shot</code> format, checkpoint-328).
QLoRA+ExtraData	Same QLoRA recipe trained on the Cochrane train split augmented with external 1→1 rephrase pairs (PLABA + Med-EASi, +8,587 pairs; checkpoint-433).

Qwen3.5-2B: SARI and BI ■ SARI ■ BLEU paraphrase test split (N=667)



Qwen3.5-2B: BERTScore on Cochrane rephrase test split (N=667)

